

KSQ L DB

Stream processing Made Simple with Kafka

ksqlDB

The database purpose-built for stream processing applications.

build stream processing applications on top of Apache Kafka with the ease of building traditional applications on a relational database using SQL



Stream Processing in Action - KSQL

data stream emitted by a production line

```
{  
  "reading_ts": "2021-06-24T09:30:00-05:00",  
  "sensor_id": "aa-101",  
  "production_line": "w01",  
  "widget_type": "acme94",  
  "temp_celcius": 23,  
  "widget_weight_g": 100  
}
```

- Alert if the line produces an item that is over a threshold weight
- Monitor the rate of item production
- Detect anomalies in the equipment
- Store the data for analytics dashboards and ad hoc querying

```
SELECT *  
FROM WIDGETS  
WHERE WEIGHT_G > 120
```

```
SELECT COUNT(*)  
FROM WIDGETS  
WINDOW TUMBLING (SIZE 1 HOUR)  
GROUP BY PRODUCTION_LINE
```

```
SELECT AVG(TEMP_CELCIUS) AS TEMP  
FROM WIDGETS  
WINDOW TUMBLING (SIZE 5 MINUTES)  
GROUP BY SENSOR_ID  
HAVING TEMP > 20
```

```
CREATE SINK CONNECTOR dw WITH (  
  connector.class = S3Connector,  
  topics = widgets  
  [...]  
);
```

How does ksqlDB work?

ksqlDB



Compute

Kafka




Storage



- ksqlDB allows us to read, filter, transform, or otherwise process streams and tables of events, which are backed by Kafka topics.
- We can also join streams and/or tables to meet the needs of our application.
- And we can do all of this using familiar SQL syntax.

Interacting with ksqlDB

Command Line Interface (CLI)



Copyright 2017-2021 Confluent Inc.

CLI v0.18.0, Server v0.18.0 located at http://ksqldb:8088
Server Status: RUNNING

Having trouble? Type 'help' (case-insensitive) for a rundown of how things work!

ksql> █

Web UI

Editor Flow Streams Tables Running queries

```
1 CREATE OR REPLACE STREAM SHIP_STATUS_REPORTS WITH (KAFKA_TOPIC='SHIP_STATUS_REPORTS', PAR
2 STATUS_REPORT.ROWTIME STATUS_TIMESTAMP,
3 STATUS_REPORT.*,
4 SHIP_INFO.*,
5 STRUCT('lat':=STATUS_REPORT.LAT, 'lon':=STATUS_REPORT.LON) LOCATION
6 FROM AIS_MSG_TYPE_1_2_3 STATUS_REPORT
7 LEFT OUTER JOIN SHIP_INFO SHIP_INFO ON ((STATUS_REPORT.MMSI = SHIP_INFO.MMSI))
8 EMIT CHANGES;
```

● [Add query properties](#)

Stop

Run query

REST API



Using ksqlDB

create a **stream** called MOVEMENTS

```
CREATE STREAM MOVEMENTS (PERSON VARCHAR KEY, LOCATION VARCHAR) WITH  
(VALUE_FORMAT='JSON', PARTITIONS=1, KAFKA_TOPIC='movements');
```

```
INSERT INTO MOVEMENTS VALUES ('Allison', 'Denver');  
INSERT INTO MOVEMENTS VALUES ('Robin', 'Leeds');  
INSERT INTO MOVEMENTS VALUES ('Robin', 'Ilkley');  
INSERT INTO MOVEMENTS VALUES ('Allison', 'Boulder');
```

Using ksqlDB

```
#show streams;
```

```
#SET 'auto.offset.reset' = 'earliest';
```

(to see all of the data that's already in the stream, and not just new messages as they arrive)

Now run the following SELECT to show all of the events in MOVEMENTS:

```
#SELECT * FROM MOVEMENTS EMIT CHANGES;
```


Using ksqlDB

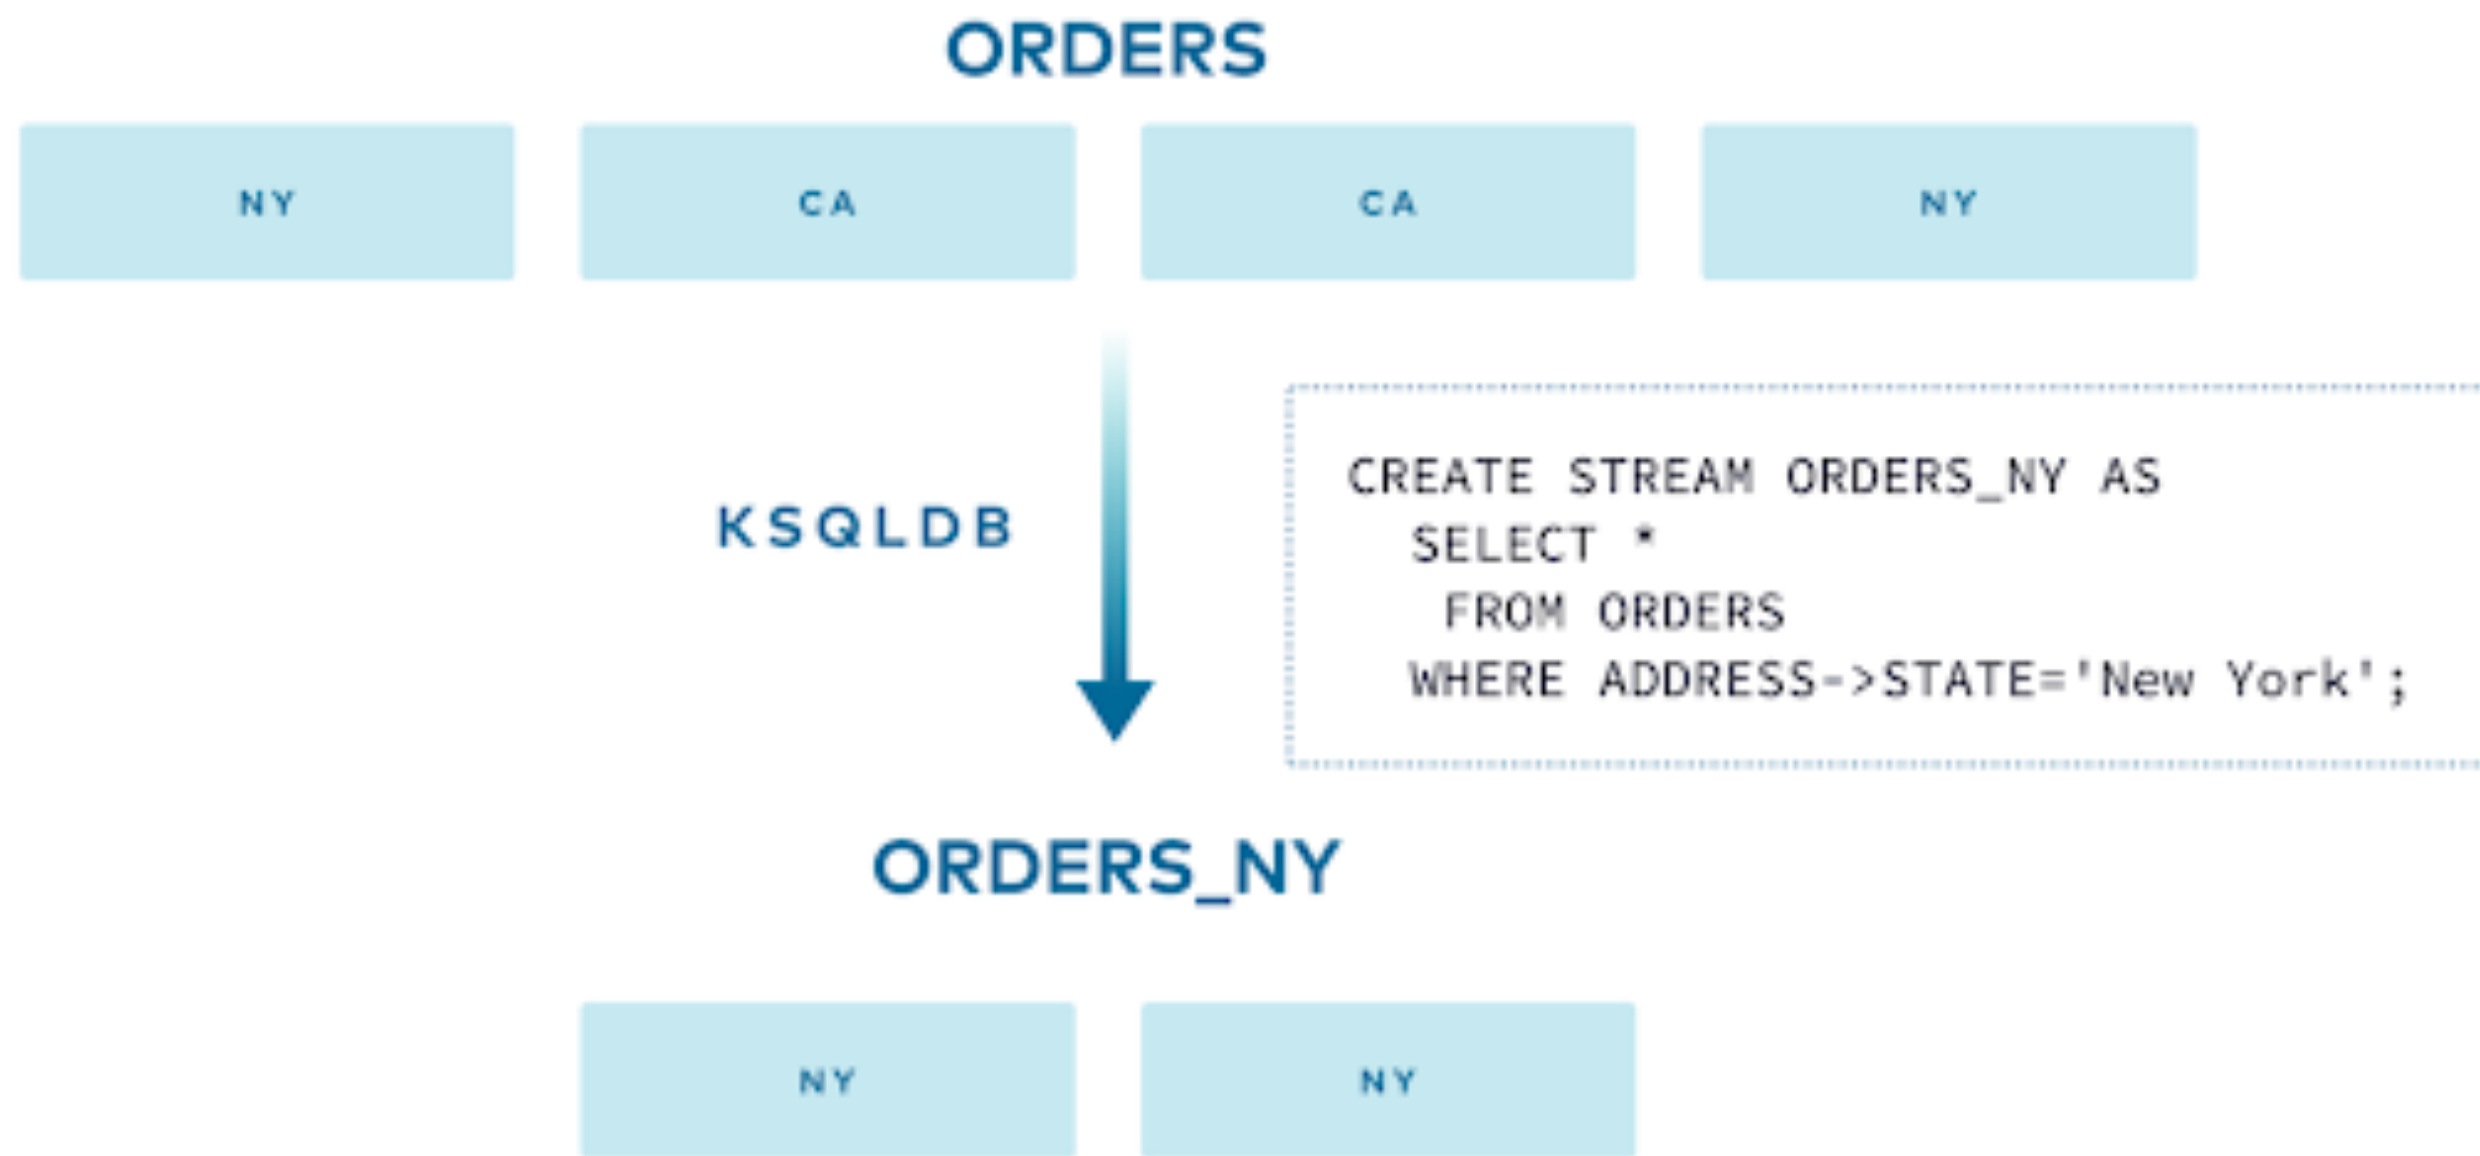
```
CREATE TABLE PERSON_STATS WITH (VALUE_FORMAT='AVRO') AS
SELECT PERSON,
       LATEST_BY_OFFSET(LOCATION) AS LATEST_LOCATION,
       COUNT(*) AS LOCATION_CHANGES,
       COUNT_DISTINCT(LOCATION) AS UNIQUE_LOCATIONS
FROM MOVEMENTS
GROUP BY PERSON
EMIT CHANGES;
```

#show tables;

This will verify that the table has been created.

```
SELECT * FROM PERSON_STATS WHERE person ='Allison';
```


Filtering with ksqlDB



ksqlDB streams are backed by Kafka topics, this means that you are writing this data Directly to a Kafka topic:

```
ksql> SHOW TOPICS;
```

Kafka Topic	Partitions	Partition Replicas
ORDERS_NY	6	1
orders	6	1

```
ksql> PRINT ORDERS_NY LIMIT 2;  
Key format: ^\_(\ツ)\_/^ - no data processed  
Value format: AVRO  
rowtime: 2021/02/22 11:16:10.881 Z, key: <null>, value: {"ORDERTIME": 1613851200000, "ORDERID": 1, "ADDRESS": {"STATE": "New York", "CITY": "New York"}, "AMOUNT": 100, "CUSTID": 1, "CUSTNAME": "John Doe"},  
rowtime: 2021/02/22 10:57:36.879 Z, key: <null>, value: {"ORDERTIME": 1613851200000, "ORDERID": 1, "ADDRESS": {"STATE": "New York", "CITY": "New York"}, "AMOUNT": 100, "CUSTID": 1, "CUSTNAME": "John Doe"},  
Topic printing ceased  
ksql>
```

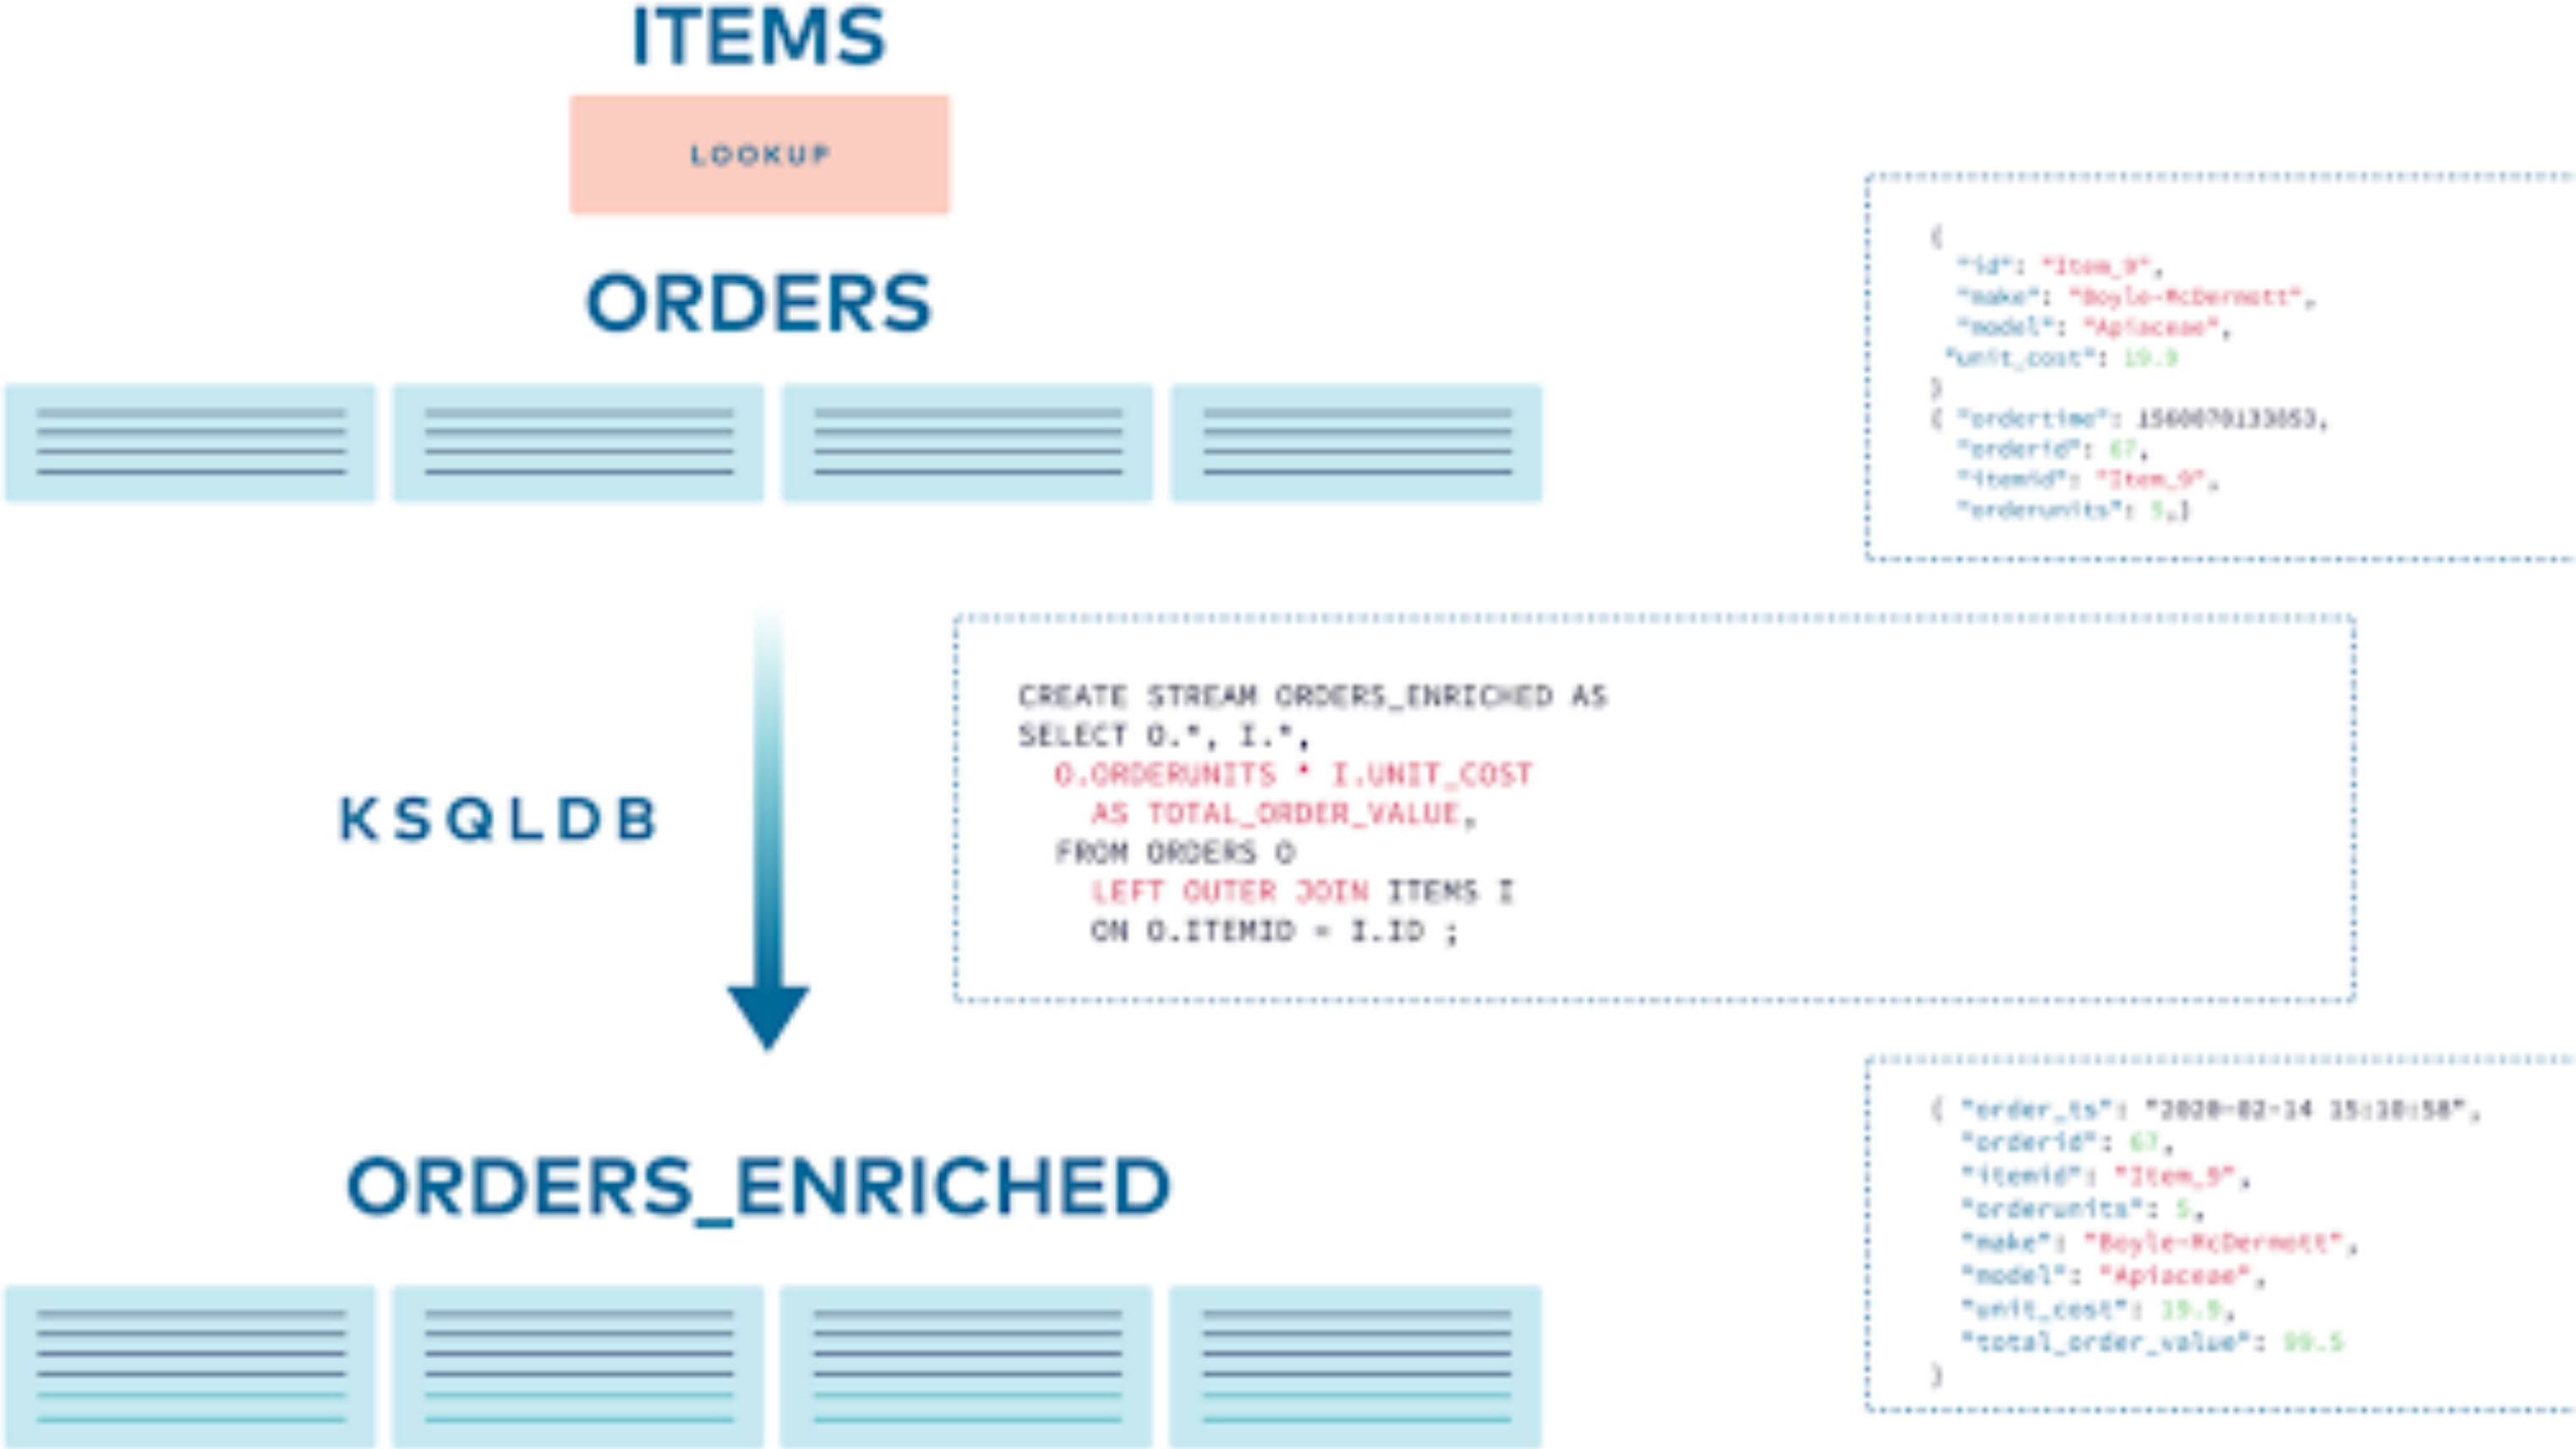
Lookups and Joins with ksqlDB

Joins can be between streams of events, between streams and tables, or between tables and tables.

Instead of making calls out to external systems to enrich the data with lookups (which introduces complexity, and can be inefficient), the data is ingested, using Kafka Connect, from the external system into an Apache Kafka topic, from which the join can be performed natively using ksqlDB.



Lookups and Joins with ksqlDB



Join - Example

joining order events to a table that contains information about the item being ordered:

```
CREATE STREAM ORDERS_ENRICHED AS
```

```
SELECT
```

```
O.*, I.*, O.ORDERUNITS * I.UNIT_COST AS TOTAL_ORDER_VALUE
```

```
FROM ORDERS O
```

```
LEFT OUTER JOIN ITEMS I
```

```
ON O.ITEMID = I.ID ;
```

Converting Data Formats with ksqlDB

There are different ways to serialize data written to Apache Kafka topics. Common options include Avro, Protobuf, and JSON.

You can use ksqlDB to create a new stream of data identical to the source but serialized differently.

To write a stream of data from its CSV source to a stream using Protobuf, you would first declare the schema of the CSV data:

```
CREATE STREAM source_csv_stream (ITEM_ID INT,  
                                DESCRIPTION VARCHAR,  
                                UNIT_COST DOUBLE,  
                                COLOUR VARCHAR,  
                                HEIGHT_CM INT,  
                                WIDTH_CM INT,  
                                DEPTH_CM INT)  
WITH (KAFKA_TOPIC = 'source_topic',  
      VALUE_FORMAT = 'DELIMITED');
```

```
CREATE STREAM target_proto_stream  
WITH (VALUE_FORMAT = 'PROTOBUF')  
AS SELECT * FROM source_csv_stream
```

Push Queries and Pull Queries

Push queries are identified by the `EMIT CHANGES` clause. By running a push query, the client will receive a message for every change that occurs on the stream (that is, every new message).

Pull queries return the current state to the client, and then terminate. In that sense, they are much more akin to a `SELECT` statement executed on a regular RDBMS. They can only be used against `ksqlDB` tables with materialized state, that is, a table in which there is an aggregate function. Currently, tables declared against an existing Apache Kafka topic cannot be queried with a pull query (as of `ksqlDB v0.15` / February 2021).

Lab : **Workflow using KSQL - CLI – 90 Minutes**